

✓ Assignment 5

This is an basecode for assignment 5 of Artificial Intelligence class (CSCE-4613), Spring 2025

```
import torch
import torchvision
import torch.nn as nn
from torchvision import transforms

from PIL import Image
import pickle
import matplotlib.pyplot as plt
```

✓ Question 1

Define Input Transformation

```
MEAN = [0.485, 0.456, 0.406]
STD = [0.229, 0.224, 0.225]
IMAGE_SIZE = 224
image_transforms = transforms.Compose([
    transforms.Resize(IMAGE_SIZE),
    transforms.CenterCrop(IMAGE_SIZE),
    transforms.ToTensor(),
    transforms.Normalize(mean = MEAN, std = STD)])

index2name = pickle.load(open("imagenet_class_names.pkl", "rb"))
```

✓ Define Model

```
model = torchvision.models.resnet50(pretrained=True)
softmax_layer = nn.Softmax(dim=1)
model.eval()
```



```

        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False,
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (4): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (5): Bottleneck(
        (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

✓ Classify and Visualize Image

```

image_path = "dog.jpg"
original_image = Image.open(image_path).convert("RGB")
image = image_transforms(original_image)
image = image.unsqueeze(0)
output = softmax_layer(model(image))

prediction = torch.argmax(output, dim=1).item()
prob = output[0, prediction].item() * 100
predicted_name = index2name[prediction]

plt.imshow(original_image)
plt.title("Class: %s. Probabilty: %.2f" % (predicted_name, prob) + "%")
plt.axis("off")
plt.show()

```



Class: golden retriever. Probabilty: 93.91%



✓ Get Top-K Predictions

```

K = 5
# Write your code to get the top-K predictions (top K classes that have highest prob)
top_k_values, top_k_indices = torch.topk(output, K, dim=1)
top_k_probs = top_k_values[0].tolist()
top_k_classes = top_k_indices[0].tolist()

# Print top-K predictions
print("Top {} predictions:".format(K))
for i in range(K):
    class_name = index2name[top_k_classes[i]]
    probability = top_k_probs[i] * 100
    print("{}{}. {} - {:.2f}%".format(i+1, class_name, probability))

```



Top 5 predictions:

1. golden retriever - 93.91%
2. Labrador retriever - 3.65%
3. Brittany spaniel - 0.69%

- 4. clumber, clumber spaniel - 0.30%
- 5. tennis ball - 0.29%

✓ Question 2

Define Training Data Loader

```
train_batch_size = 32
train_transforms = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomCrop(32, padding=4),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2470, 0.2435, 0.2616))
])

# Define training dataset and data loader
train_dataset = torchvision.datasets.CIFAR10(root='data/CIFAR-10', train=True,
                                              download=True, transform=train_transforms)
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=train_batch_size,
                                           shuffle=True, num_workers=2)
```

➡ 100%|██████████| 170M/170M [00:03<00:00, 48.3MB/s]

✓ Define Model and Training Framework

```

cuda = torch.cuda.is_available()
# model = None # Define a classification model for 10 classes
# Define a simple CNN for CIFAR-10
class CIFAR10CNN(nn.Module):
    def __init__(self):
        super(CIFAR10CNN, self).__init__()
        # Convolutional layers
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)

        # Pooling layer
        self.pool = nn.MaxPool2d(2, 2)

        # Fully connected layers
        self.fc1 = nn.Linear(128 * 4 * 4, 512)
        self.fc2 = nn.Linear(512, 10)

        # Dropout
        self.dropout = nn.Dropout(0.2)

        # Batch normalization
        self.bn1 = nn.BatchNorm2d(32)
        self.bn2 = nn.BatchNorm2d(64)
        self.bn3 = nn.BatchNorm2d(128)

    def forward(self, x):
        # Apply convolution, batch norm, ReLU, and pooling
        x = self.pool(torch.relu(self.bn1(self.conv1(x))))
        x = self.pool(torch.relu(self.bn2(self.conv2(x))))
        x = self.pool(torch.relu(self.bn3(self.conv3(x))))

        # Flatten
        x = x.view(-1, 128 * 4 * 4)

        # Fully connected with dropout
        x = self.dropout(torch.relu(self.fc1(x)))
        x = self.fc2(x)

        return x

# Initialize the model
model = CIFAR10CNN()

if cuda:
    model.cuda()

model.train()
learning_rate = 0.001
num_epochs = 3
optim = torch.optim.Adam(model.parameters(), lr = learning_rate)
loss_fn = nn.CrossEntropyLoss()

```



```

loss_logger = []
accuracy_logger = []

for epoch in range(1, num_epochs + 1):
    for it, (images, labels) in enumerate(train_loader):
        if cuda:
            images = images.cuda()
            labels = labels.cuda()

            # Write your code to computer outputs, loss, accuracy, and optimize model
            # outptus = ???
            # loss = ???
            # accuracy = ???
            # optimize the model
            # loss = None
            # accuracy = None
            # Forward pass
            outputs = model(images)
            loss = loss_fn(outputs, labels)

            # Backward and optimize
            optim.zero_grad()
            loss.backward()
            optim.step()

            # Calculate accuracy
            _, predicted = torch.max(outputs.data, 1)
            total = labels.size(0)
            correct = (predicted == labels).sum().item()
            accuracy = correct / total * 100

            loss_logger.append(loss.item())
            accuracy_logger.append(accuracy)

        if it % 200 == 0:
            print("Epoch [%d/%d]. Iter [%d/%d]. Loss: %.4f. Accuracy: %.2f" % (epoch, num_epochs,
torch.save(model.state_dict(), "CIFAR10-Model.pth")

# Write your code to visualize the training losses over interations
plt.figure(figsize=(10, 5))
plt.plot(loss_logger)
plt.title('Training Loss Over Iterations')
plt.xlabel('Iteration')
plt.ylabel('Loss')
plt.grid(True)
plt.show()

# Write your code to visualize the training accuracies over interations
plt.figure(figsize=(10, 5))
plt.plot(accuracy_logger)
plt.title('Training Accuracy Over Iterations')


```

```
plt.xlabel('Iteration')  
plt.ylabel('Accuracy (%)')  
plt.grid(True)  
plt.show()
```





```
Epoch [1/3]. Iter [1/1563]. Loss: 2.2916. Accuracy: 12.50
Epoch [1/3]. Iter [201/1563]. Loss: 1.4965. Accuracy: 40.62
Epoch [1/3]. Iter [401/1563]. Loss: 1.5610. Accuracy: 46.88
Epoch [1/3]. Iter [601/1563]. Loss: 1.4966. Accuracy: 40.62
Epoch [1/3]. Iter [801/1563]. Loss: 1.4353. Accuracy: 43.75
Epoch [1/3]. Iter [1001/1563]. Loss: 1.1127. Accuracy: 62.50
Epoch [1/3]. Iter [1201/1563]. Loss: 1.2591. Accuracy: 56.25
Epoch [1/3]. Iter [1401/1563]. Loss: 1.1439. Accuracy: 56.25
Epoch [2/3]. Iter [1/1563]. Loss: 1.4925. Accuracy: 46.88
Epoch [2/3]. Iter [201/1563]. Loss: 0.9655. Accuracy: 59.38
Epoch [2/3]. Iter [401/1563]. Loss: 0.8149. Accuracy: 71.88
Epoch [2/3]. Iter [601/1563]. Loss: 1.0035. Accuracy: 62.50
Epoch [2/3]. Iter [801/1563]. Loss: 1.0080. Accuracy: 65.62
Epoch [2/3]. Iter [1001/1563]. Loss: 1.1147. Accuracy: 62.50
Epoch [2/3]. Iter [1201/1563]. Loss: 1.0802. Accuracy: 65.62
Epoch [2/3]. Iter [1401/1563]. Loss: 0.7983. Accuracy: 68.75
Epoch [3/3]. Iter [1/1563]. Loss: 0.9694. Accuracy: 65.62
Epoch [3/3]. Iter [201/1563]. Loss: 1.0216. Accuracy: 62.50
Epoch [3/3]. Iter [401/1563]. Loss: 0.7440. Accuracy: 71.88
Epoch [3/3]. Iter [601/1563]. Loss: 0.8917. Accuracy: 62.50
Epoch [3/3]. Iter [801/1563]. Loss: 0.8521. Accuracy: 71.88
Epoch [3/3]. Iter [1001/1563]. Loss: 1.3275. Accuracy: 59.38
Epoch [3/3]. Iter [1201/1563]. Loss: 0.7498. Accuracy: 81.25
Epoch [3/3]. Iter [1401/1563]. Loss: 0.7952. Accuracy: 71.88
```

Training Loss Over Iterations



Training Accuracy Over Iterations





✓ Load Model and Evaluate Model On Testing Dataset

```

cuda = torch.cuda.is_available()
# model = # Define a classification model for 10 classes
model = CIFAR10CNN()
if cuda:
    model.cuda()

# Load your trained model in the previous step
model.load_state_dict(torch.load("CIFAR10-Model.pth"))
model.eval()

# test_dataset = None # Define Testing Set of CIFAR-10
# test_loader = None # Define Testing Loader of CIFAR-10
# Define test dataset and loader
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2470, 0.2435, 0.2616))
])

test_dataset = torchvision.datasets.CIFAR10(root='data/CIFAR-10', train=False,
                                             download=True, transform=test_transforms)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=100,
                                           shuffle=False, num_workers=2)

all_predictions = []
all_labels = []

final_accuracy = 0.0
final_top_5_acuuracy = 0.0

for it, (images, labels) in enumerate(test_loader):
    if cuda:
        images = images.cuda()
        labels = labels.cuda()

    # Write your code to computer outputs, accuracy, and top 5 accuracy
    # outptus = ???
    # accuracy = ???
    # top_5_acuuracy = ???
    # accuracy = None
    # top_5_acuuracy = None
    # Forward pass

```

```
outputs = model(images)

# Calculate accuracy
_, predicted = torch.max(outputs.data, 1)
total = labels.size(0)
correct = (predicted == labels).sum().item()
accuracy = correct / total * 100

# Calculate top-5 accuracy
_, top5_indices = torch.topk(outputs, 5, dim=1)
batch_correct = 0
for i in range(len(labels)):
    if labels[i] in top5_indices[i]:
        batch_correct += 1
top_5_acuuracy = batch_correct / total * 100
all_predictions.extend(predicted.cpu().numpy())
all_labels.extend(labels.cpu().numpy())
final_accuracy += accuracy
final_top_5_acuuracy += top_5_acuuracy

if it % 500 == 0:
    print("Iter [%d/%d]. Accuracy: %.2f" % (it + 1, len(test_loader), accuracy))

print("Final Accuracy: %0.2f" % (final_accuracy / len(test_loader)))
print("Top 5 Accuracy: %0.3f" % (top_5_acuuracy / len(test_loader)))

from sklearn.metrics import confusion_matrix
import numpy as np
import seaborn as sns

# CIFAR-10 class names
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'hor
cm = confusion_matrix(all_labels, all_predictions)
```